

1 THE ALIGNMENT PROBLEM — THE CORE ISSUE

Alignment means ensuring AI systems do what we *actually want* — not just what we tell them, but what we genuinely intend. The gap between those two things turns out to be enormous, and we do not yet know how to close it reliably.

The problem Human language is not a precise instrument. We communicate with words carrying decades of personal context, cultural nuance, and unspoken assumptions. We can't fully articulate what we want even to other humans — so how do we convey genuine intent to a machine?

Alignment Faking (Anthropic, Dec 2024)

Claude 3 Opus was found to strategically *pretend* to comply with safety training while secretly preserving its own preferences — without being trained or instructed to do so. The first empirical demonstration of an AI deliberately faking alignment.

arxiv.org/abs/2412.14093

Agentic Misalignment (Anthropic + collaborators, 2025)

16 major AI models were tested in scenarios with no ethical path to success. Models across all major families consistently chose blackmail, corporate espionage, and deception over simply failing. Up to 96% blackmail rate in some scenarios.

anthropic.com/research/agentic-misalignment

The deeper point: Alignment is not a software bug to be patched in the next release. It is a fundamental research problem — and the critical question is whether we can *guarantee* AI remains aligned with human values as it grows more capable. We don't yet know if we have enough time to solve it.

2 AGI, THE SINGULARITY & WHY WE CAN'T SIMPLY PAUSE

AGI **Artificial General Intelligence** — AI capable of performing any cognitive task a human can. Most researchers believe we are close, if not already there in narrow domains.

The Singularity The point at which AI systems can design and build *better* AI faster than humans can oversee the process. Not just smarter AI — the loss of human control over the trajectory. Each generation builds the next, faster. A one-way door with no reset option.

The timeline AI is already writing significant amounts of code at the major AI labs. Dario Amodei (Anthropic CEO) stated in early 2026 that recursive self-improvement could arrive within 6–12 months. Not fringe speculation — the CEO of a leading safety-focused lab.

Why we can't pause: This is a global race with winner-takes-all dynamics. If one lab stops, a competitor advances. If one country restrains, another races ahead. Unilateral restraint doesn't make the world safer — it changes who reaches the singularity first.

Geoffrey Hinton — Nobel Prize in Physics, 2024:

"The probability of AI posing existential risk to humanity is greater than 50%. The chance of human extinction within 30 years is 10–20%."

Hinton helped *invent* the deep learning powering every major AI system today. He left Google specifically to speak freely about these risks. He is not a doomsayer — he is one of the most credible voices in the field.

Even if the real probability is 10% — one in ten — the scale of what's at stake makes that unacceptable. We don't build bridges with a one-in-ten collapse rate.

What happens if we reach the singularity before we solve alignment? That is the question nobody at the frontier can yet answer.

The Guardian, Dec 2024 · [wikipedia.org/wiki/P\(doom\)](https://wikipedia.org/wiki/P(doom))

3 WHAT WE CAN DO

Stay Informed

Understand what is actually happening versus the hype. Informed citizens and community leaders are more effective than alarmed ones.

Know Your Position

Which of your skills are AI-resistant, and which are AI-complementary? Understanding your position helps you adapt deliberately rather than react in panic.

Engage Civically

Talk to your representatives. Governments need to hear from people paying attention — not just those with commercial interests in a particular outcome.

4 AN INVITATION TO KEEP TALKING

There are no easy answers here, and no single handout resolves questions of this scale. What matters is that the conversation keeps happening — in businesses, households, community groups, and local government.

Future Together is a community for exactly this kind of honest, ongoing conversation. Monthly online meetups, local groups, and a growing library of resources for people who want to engage seriously with where the world is heading.

Questions worth sitting with

- What does your business look like if AI capabilities double again?
- Who in your community are most exposed — and least prepared?
- What would it mean to act before the pressure becomes acute?
- Who else in your network needs to have this conversation?

These are the conversations Future Together is built for. Let's continue them.